



Multi Threading

Krishantha Dinesh
@ 2013-Oct

2000 West Park Drive
Westborough MA 01581 USA
Phone: 508 389 7300 Fax: 508 366 9901

The entire contents of this document are subject to copyright with all rights reserved. All copyrightable text and graphics, the selection, arrangement and presentation of all information and the overall design of the document are the sole and exclusive property of Virtusa.

Copyright © 2012 Virtusa Corporation. All rights reserved

Objectives

- Not to teach Java

Prerequisites

- `Java.basicknowledge >(60/100);`
- `Mobilephone.soundmode=soundmode.completesilent && soundmode.completesilent != soundmode.vibrate`

Objectives

- To understand the purpose of multithreading
- To describe Java's multithreading mechanism
- To explain concurrency issues caused by multithreading
- To outline synchronized access to shared resources

What is Multithreading?

- Multithreading is similar to multi-processing.
- A multi-processing Operating System can run several processes at the same time and each process has its own address/memory space
- The OS's scheduler decides when each process is executed
- Only one process is actually executing at any given time. However, the system appears to be running several programs simultaneously
- Separate processes do not have access to each other's memory space
- In a multithreaded application, there are several points of execution within the same memory space.
- Each point of execution is called a thread
- Threads share access to memory

Why use multithreading

- In a single threaded application, one thread of execution must do everything
- If an application has several tasks to perform, those tasks will be performed when the thread can get to them.
- A single task which requires a lot of processing can make the entire application appear to be "sluggish" or unresponsive.
- In a multithreaded application, each task can be performed by a separate thread
- If one thread is executing a long process, it does not make the entire application wait for it to finish.
- If a multithreaded application is being executed on a system that has multiple processors, the OS may execute separate threads simultaneously on separate processors.

How it is worked

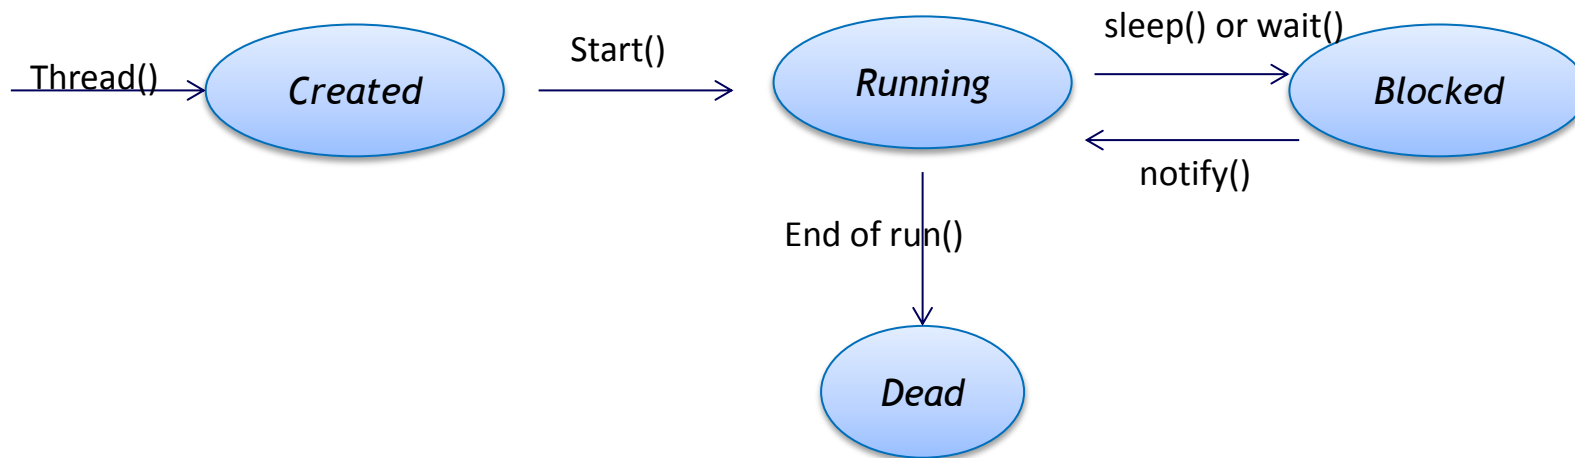
- Each thread is given its own "context"
- A thread's context includes virtual registers and its own calling stack
- The "scheduler" decides which thread executes at any given time
- Since many OS now directly support multithreading, the VM may use the system's scheduler for scheduling threads
- The scheduler maintains a list of ready threads (the run queue) and a list of threads waiting for input (the wait queue)
- Each thread has a priority. The scheduler typically schedules between the highest priority threads in the run queue
- Note: the programmer cannot make assumptions about how threads are going to be scheduled. Typically, threads will be executed differently on different platforms.

Multithreading in java

- Few programming languages directly support threading
- Although many have add-on thread support
- Add on thread support is often quite cumbersome to use
- The Java Virtual machine has its own runtime threads
- Threads are represented by a Thread class
- A thread object maintains the state of the thread
- It provides control methods such as interrupt, start, sleep, yield, wait
- When an application executes, the main method is executed by a single thread.
- If the application requires more threads, the application must create them.

States

- Threads can be in one of four states
 - Created, Running, Blocked, and Dead
- A thread's state changes based on:
 - Control methods such as start, sleep, yield, wait, notify
- Termination of the run method



How does thread run


- The thread class has a run() method
- run() is executed when the thread's start() method is invoked
- The thread terminates if the run method terminates
- To prevent a thread from terminating, the run method must not end
- run methods often have an endless loop to prevent thread termination
- One thread starts another by calling its start method
- The sequence of events can be confusing to those more familiar with a single threaded model.

Create own threads

- The obvious way to create your own threads is to subclass the Thread class and then override the run() method
- This is the easiest way to do it also It is not the recommended way to do it.
- Because threads are usually associated with a task, the object which provides the run method is usually a subclass of some other class
- If it inherits from another class, it cannot inherit from Thread.
- The solution is provided by an interface called Runnable.
- Runnable defines one method - public void run()
- One of the Thread classes constructor takes a reference to a Runnable object
- When the thread is started, it invokes the run method in the runnable object instead of its own run method.

Runnable

- In the example below, when the Thread object is instantiated, it is passed a reference to a "Runnable" object
- The Runnable object must implement a method called "run"
- When the thread object receives a start message, it checks to see if it has a reference to a Runnable object:
- If it does, it runs the "run" method of that object
- If not, it runs its own "run" method



example

```
package com.virtusa.training.java.thread;

public class ThreadSample implements Runnable {

    private Thread thisThread;

    public void start() {

        if (thisThread == null) {
            thisThread = new Thread();
            thisThread.start();
        }
    }

    @Override
    public void run() {
        // this run on its own thread
    }
}
```

Termination

- In Java 1.1, the Thread class had a stop() method
- One thread could terminate another by invoking its stop() method.
- However, using stop() could lead to deadlocks
- The stop() method is now deprecated. DO NOT use the stop method to terminate a thread
- The correct way to stop a thread is to have the run method terminate
- Add a boolean variable which indicates whether the thread should continue or not
- Provide a set method for that variable which can be invoked by another thread

Stop - example

```
package com.virtusa.training.java.thread;

public class TerminationSample implements Runnable {

    private Thread thisThread;
    private boolean isToStop;

    public void setToStop(boolean isToStop) {
        this.isToStop = isToStop;
    }

    public void start() {

        if (thisThread == null) {
            thisThread = new Thread();
            thisThread.start();
        }

    }

    @Override
    public void run() {
        while (true) {
            if (isToStop)
                break;
        }
    }
}
```

Create multiple Threads

- The previous example illustrates a Runnable class which creates its own thread when the start method is invoked.
- If one wished to create multiple threads, one could simple create multiple instances of the Runnable class and send each object a start message
- Each instance would create its own thread object
- Is the a maximum number of threads which can be created?
- There is no defined maximum in Java.
- If the VM is delegating threads to the OS, then this is platform dependent.
- A good rule of thumb for maximum thread count is to allow 2Mb of ram for each thread
- Although threads share the same memory space, this can be a reasonable estimate of how many threads your machine can handle.

Priorities

- Every thread is assigned a priority (between 1 and 10) The default is 5
- The higher the number, the higher the priority Can be set with `setPriority`
- The standard mode of operation is that the scheduler executes threads with higher priorities first.
- This simple scheduling algorithm can cause problems. Specifically, one high priority thread can become a "CPU hog".
- A thread using vast amounts of CPU can share CPU time with other threads by invoking the `yield()` method on itself.
- The more CPU a thread receives, the lower its priority becomes
- The more a thread waits for the CPU, the higher its priority becomes
- Because of thread aging, the effect of setting a thread's priority is dependent on the platform

Yield() and sleep()

- Sometimes a thread can determine that it has nothing to do
- Sometimes the system can determine this. ie. waiting for I/O
- When a thread has nothing to do, it should not use CPU
- This is called a busy-wait.
- Threads in busy-wait are busy using up the CPU doing nothing.
- Often, threads in busy-wait are continually checking a flag to see if there is anything to do.
- It is worthwhile to run a CPU monitor program on your desktop
- You can see that a thread is in busy-wait when the CPU monitor goes up (usually to 100%), but the application doesn't seem to be doing anything.

Cont..

- Threads in busy-wait should be moved from the Run queue to the Wait queue so that they do not hog the CPU
- Use `yield()` or `sleep(time)`
- Yield simply tells the scheduler to schedule another thread
- Sleep guarantees that this thread will remain in the wait queue for the specified number of milliseconds.

Concurrent access

- Those familiar with databases will understand that concurrent access to data can lead to data integrity problems
- Specifically, if two sources attempt to update the same data at the same time, the result of the data can be undefined.
- The outcome is determined by how the scheduler schedules the two sources.
- Since the schedulers activities cannot be predicted, the outcome cannot be predicted
- Databases deal with this mechanism through "locking"
- If a source is going to update a table or record, it can lock the table or record until such time that the data has been successfully updated.
- While locked, all access is blocked except to the source which holds the lock.
- Java has the equivalent mechanism. It is called synchronization

Synchronization

- To obtain the lock, you must synchronize with the object
- The simplest way to use synchronization is by declaring one or more methods to be synchronized
- When a synchronized method is invoked, the calling thread attempts to obtain the lock on the object.
- if it cannot obtain the lock, the thread goes to sleep until the lock becomes available
- Once the lock is obtained, no other thread can obtain the lock until it is released. ie, the synchronized method terminates
- When a thread is within a synchronized method, it knows that no other synchronized method can be invoked by any other thread
- Therefore, it is within synchronized methods that critical data is updated

Provide thread safe

- If an object contains data which may be updated from multiple thread sources, the object should be implemented in a thread-safe manner
- All access to critical data should only be provided through synchronized methods (or synchronized blocks).
- In this way, we are guaranteed that the data will be updated by only one thread at a time.

```
package com.virtusa.training.java.thread;

public class ThreadSafe {

    private double balance;

    public synchronized void withdraw(float amount)
    {
        if ((amount>0.0) && (amount<=balance))
            balance = balance - amount;
    }

    public synchronized void deposit(float anAmount)
    {
        if (anAmount>0.0)
            balance = balance + anAmount;
    }
}
```

Consequences

- Many threads may be waiting to gain access to one of the object's synchronized methods
- The object remains locked as long as a thread is within a synchronized method.
- Ideally, the method should be kept as short as possible.
- Another solution is to provide synchronization on a block of code instead of the entire method
- In this case, the object's lock is only held for the time that the thread is within the block.
- The intent is that we only lock the region of code which requires access to the critical data. Any other code within the method can occur without the lock.
- In high load situations where multiple threads are attempting to access critical data, this is by far a much better implementation

solution

```
package com.virtusa.training.java.thread;

public class ThreadSafe {

    private double balance;

    public void withdraw(float amount) {

        synchronized (this) {
            if (amount <= balance)
                balance = balance - amount;
        }
    }

    public void deposit(float anAmount) {

        synchronized (this) {
            balance = balance + anAmount;
        }
    }
}
```


Thread sleep example

```
package com.virtusa.training.java.thread.sleep.D6;

public class ThreadClass implements Runnable {

    @Override
    public void run() {
        while (true) {
            System.out.println(Thread.currentThread().getName());
            try {
                Thread.sleep(100);
            } catch (InterruptedException e) {
                e.printStackTrace();
            }
        }
    }
}

package com.virtusa.training.java.thread.sleep.D6;

public class ThreadRunner {

    public static void main(String[] args) {
        ThreadClass fromThread = new ThreadClass();

        new Thread(fromThread, "Nala").start();

        new Thread(fromThread, "Damayanthi").start();
    }
}
```

Yield example

```
package com.virtusa.training.java.thread.yield.D5;

public class ThreadClass implements Runnable {

    @Override
    public void run() {
        while (true) {
            System.out.println(Thread.currentThread().getName());
            Thread.yield();
        }
    }
}

package com.virtusa.training.java.thread.yield.D5;

public class ThreadRunner {

    public static void main(String[] args) {
        ThreadClass fromThread = new ThreadClass();

        new Thread(fromThread, "Nala").start();

        new Thread(fromThread, "Damayanthi").start();
    }
}
```

Termination example

```
package com.virtusa.training.thread.terminate;

public class ThreadTermination implements Runnable {

    private boolean isStop;

    @Override
    public void run() {
        while (!isStop) {
            System.out.println("i am a live");
        }
        System.out.println("oops i am dead");
    }

    public void setStop(boolean isStop) {
        this.isStop = isStop;
    }
}
```

```
package com.virtusa.training.thread.terminate;

public class ThreadRunner {

    public static void main(String[] args) {
        ThreadTermination fromThread = new ThreadTermination();
        Thread thread = new Thread(fromThread, "Nala");
        thread.start();
        for (int i = 0; i < 1000000; i++) {

            if (i == 50000) {
                fromThread.setStop(true);
            }
        }
    }
}
```

Yield tester

```
package com.virtusa.training.thread.yield;

public class YieldSampleClass implements Runnable {

    private static int Ganga,yamuna;
    private String name;
    @Override
    public void run() {

        while (true) {
            name=(Thread.currentThread().getName());

            if (name.equals("Ganga"))
                Ganga++;
            else
                yamuna++;

            System.out.println("G: "+Ganga + "="+"Y: "+yamuna);
            Thread.yield();
        }
    }
}
```

```
package com.virtusa.training.thread.yield;

public class ThreadRunner {

    public static void main(String[] args) {
        YieldSampleClass fromThread = new YieldSampleClass();

        Thread thread = new Thread(fromThread, "Ganga");
        Thread thread2 = new Thread(fromThread, "Yamuna");

        thread.setPriority(1);
        thread2.setPriority(10);

        thread.start();
        thread2.start();
    }
}
```